

COSCS494/594 Fundamentals of Digital
Archeology
Data Discovery

Audris Mockus
University of Tennessee
audris@utk.edu

Where can I find data? (early history)

Early history

- ▶ Cultural artifacts (people's mind): 3M-5K years ago
 - ▶ Passed on verbally, through practice and joint experiences
- ▶ Physical Artifacts
 - ▶ Stone/clay tablets/papyrus/paper: 5K-now

Where can I find data? (Digital Age)

Digital Age

- ▶ Mainframe era
 - ▶ 1928 Magnetic tape
 - ▶ 1932 Magnetic Drum
 - ▶ 1956 Hard disk
- ▶ Minicomputers/PCs
 - ▶ 1980 CD, 1995 DVD
 - ▶ 1990 PCs
- ▶ 2010 Cloud era
 - ▶ SourceForge
 - ▶ Salesforce

Where can I find data? (Now)

In the cloud

- ▶ Public internet
- ▶ Corporate internet
- ▶ In what form?
 - ▶ Web pages
 - ▶ Within applications (Deep Web)

What is Deep Web

- ▶ Traditional Web Spider/Crawler
 - ▶ Start from a set of URLs
 - ▶ Retrieve and extract all the links
 - ▶ Snowball
 - ▶ follow all the new links until no more found
 - ▶ Periodically revisit all the URLs found so far to see if there are new links
- ▶ Deep Web: most of the content you can not find using the method above
 - ▶ Search: Google/BitBucket/SourceForge/GoogleCode
 - ▶ API's (typically REST), e.g, github, Twitter, Facebook
 - ▶ Other interfaces:
 - ▶ ITS: Bugzilla, Debian, Ubuntu, JIRA, ...
 - ▶ Application commands for VCS: hg, git, svn, bazaar, CVS
 - ▶ POST (web forms) interface to relational databases

REST APIs

E.g. GitHub

`https://developer.github.com/v3/`

- ▶ For example, see:

`https://api.github.com/repos/fdac/syllabus/issues?
state=closed`

- ▶ The output is json (as your iPython notebook)

```
[  
  {  
    "url": "https://api.github.com/repos/fdac/syllabus/issu  
    "labels_url": "https://api.github.com/repos/fdac/syllab  
    "comments_url": "https://api.github.com/repos/fdac/syll  
    "events_url": "https://api.github.com/repos/fdac/syllab  
  }  
  ...  
]
```

Search

E.g., GoogleCode

`http://code.google.com/hosting/search?filter=0&q=
label:SEARCHSTR&start=0`

- ▶ Page through responses
- ▶ Extract relevant links

Application commands

Extract data from VCS: list revisions

- ▶ CVS: `cvs log PRJ`
- ▶ SVN: `svn log -v --non-interactive PRJ`
- ▶ Mercurial: `hg log -v PRJ`
- ▶ Bazaar: `bzr log -v $-$long PRJ`
- ▶ GIT: `git --git-dir=PRJ log --numstat -M -C --diff-filter=ACMR
--full-history
--pretty=tformat:"%n%H;%T;%P;%an;%ae;%at;%cn;%ce;%ct;%s"`

Discovery Ethics (Don't get banned)

Don't get banned

- ▶ Check if the site has relevant policies, e.g.,
 - ▶ robots.txt
 - ▶ Text in http responses, e.g,
`https://bugzilla.gnome.org/show_bug.cgi?id=309324`
 - ▶ Rate limiting, e.g
`https://dev.twitter.com/docs/rate-limiting-faq`
- ▶ Be sensitive of potential harm you may cause
 - ▶ Bringing the server down
 - ▶ Exposing information that was not intended to be public
- ▶ Be aware of context:
 - ▶ Don't treat Google as a small project
 - ▶ Don't treat a small project as Facebook

Discovery Ethics (Workarounds)

Workarounds

- ▶ Ask for dump (e.g., Mozilla Bugzilla)
- ▶ Use authentication
 - ▶ Typically increases allowed rate
 - ▶ Sometimes provides more detail (Gnome Bugzilla)
- ▶ Create multiple accounts
- ▶ Run from multiple ip addresses
- ▶ Keep the rate below threshold

Discovery Ethics (do it once)

- ▶ Retrieve once and save

```
import pickle
r = requests .get(URL)
# to save
pickle.dump (r, open('storedReq.obj', 'w'))
# to restore
rLater = pickle.load(open('storedReq.obj', 'r'))
```

A few legal aspects

- ▶ Be careful with spiders
 - ▶ They may go to sites that you would not consider visiting
- ▶ If collecting competitive intelligence, there are legal restrictions on what you could use: make sure the obtained information is, indeed, public
- ▶ If doing a research study, you may need to pass the plan through IRRB
- ▶ There may be privacy issues:
 - ▶ Anonymize whenever possible
 - ▶ Report aggregates

A few references

- ▶ Way back machine: <http://archive.org/web/>
- ▶ http://en.wikipedia.org/wiki/Electronic_discovery
- ▶ Audris Mockus Amassing and indexing a large sample of version control systems: towards the census of public source code history In 6th IEEE Working Conference on Mining Software Repositories, May 16-17 2009. <http://mockus.org/papers/amassing-slides.pdf>
- ▶ Audris Mockus Large-scale code reuse in open source software. In ICSE'07 Intl. Workshop on Emerging Trends in FLOSS Research and Development, Minneapolis, Minnesota, May 21 2007. <http://mockus.org/papers/ossreuse.slides.pdf>

A crawler

```
while (list of unvisited URLs is not empty):  
    take URL from list  
    fetch content  
    record the content if desired  
    if content is HTML:  
        parse out URLs from links  
        foreach URL:  
            if (it matches your rules and  
                it's not already in either  
                the visited or unvisited list):  
                add it to the unvisited list
```